

2015 IEEE International Conference on Computer Vision

ICCV 2015

Table of Contents

Message from the General Chair and Program Chairs.....	xxxvi
Organizing Committee.....	xxxviii
ICCV 2015 Area Chairs.....	xl
ICCV 2015 Reviewers.....	xli
Sponsors.....	xlii

O-1A: Vision and Language

Ask Your Neurons: A Neural-Based Approach to Answering Questions about Images	1
<i>Mateusz Malinowski, Marcus Rohrbach, and Mario Fritz</i>	
Segment-Phrase Table for Semantic Segmentation, Visual Entailment and Paraphrasing	10
<i>Hamid Izadinia, Fereshteh Sadeghi, Santosh K. Divvala, Hannaneh Hajishirzi, Yejin Choi, and Ali Farhadi</i>	
Aligning Books and Movies: Towards Story-Like Visual Explanations by Watching Movies and Reading Books	19
<i>Yukun Zhu, Ryan Kiros, Rich Zemel, Ruslan Salakhutdinov, Raquel Urtasun, Antonio Torralba, and Sanja Fidler</i>	
Learning Query and Image Similarities with Ranking Canonical Correlation Analysis	28
<i>Ting Yao, Tao Mei, and Chong-Wah Ngo</i>	

P-1A: Recognition, Low-Level Vision, and Biomedical Image Analysis

Learning to See by Moving	37
<i>Pulkit Agrawal, João Carreira, and Jitendra Malik</i>	
Object Detection Using Generalization and Efficiency Balanced Co-Occurrence Features	46
<i>Haoyu Ren and Ze-Nian Li</i>	
Mining And-Or Graphs for Graph Matching and Object Discovery	55
<i>Quanshi Zhang, Ying Nian Wu, and Song-Chun Zhu</i>	

Pose Induction for Novel Object Categories	64
<i>Shubham Tulsiani, João Carreira, and Jitendra Malik</i>	
Dynamic Texture Recognition via Orthogonal Tensor Dictionary Learning	73
<i>Yuhui Quan, Yan Huang, and Hui Ji</i>	
Convolutional Channel Features	82
<i>Bin Yang, Junjie Yan, Zhen Lei, and Stan Z. Li</i>	
Local Convolutional Features with Unsupervised Training for Image Retrieval	91
<i>Mattis Paulin, Matthijs Douze, Zaid Harchaoui, Julien Mairal, Florent Perronin, and Cordelia Schmid</i>	
RIDE: Reversal Invariant Descriptor Enhancement	100
<i>Lingxi Xie, Jingdong Wang, Weiyao Lin, Bo Zhang, and Qi Tian</i>	
Discrete Tabu Search for Graph Matching	109
<i>Kamil Adamczewski, Yumin Suh, and Kyoung Mu Lee</i>	
Discriminative Learning of Deep Convolutional Feature Point Descriptors	118
<i>Edgar Simo-Serra, Eduard Trulls, Luis Ferraz, Iasonas Kokkinos, Pascal Fua, and Francesc Moreno-Noguer</i>	
Amodal Completion and Size Constancy in Natural Scenes	127
<i>Abhishek Kar, Shubham Tulsiani, João Carreira, and Jitendra Malik</i>	
Learning Where to Position Parts in 3D	136
<i>Marco Pedersoli and Tinne Tuytelaars</i>	
Query Adaptive Similarity Measure for RGB-D Object Recognition	145
<i>Yanhua Cheng, Rui Cai, Chi Zhang, Zhiwei Li, Xin Zhao, Kaiqi Huang, and Yong Rui</i>	
Listening with Your Eyes: Towards a Practical Visual Speech Recognition System	
Using Deep Boltzmann Machines	154
<i>Chao Sui, Mohammed Bennamoun, and Roberto Togneri</i>	
Cluster-Based Point Set Saliency	163
<i>Flora Ponjou Tasse, Jiri Kosinka, and Neil Dodgson</i>	
A Comprehensive Multi-Illuminant Dataset for Benchmarking of the Intrinsic Image	
Algorithms	172
<i>Shida Beigpour, Andreas Kolb, and Sven Kunz</i>	
PatchMatch-Based Automatic Lattice Detection for Near-Regular Textures	181
<i>Siying Liu, Tian-Tsong Ng, Kalyan Sunkavalli, Minh N. Do, Eli Shechtman, and Nathan Carr</i>	
A Data-Driven Metric for Comprehensive Evaluation of Saliency Models	190
<i>Jia Li, Changqun Xia, Yafei Song, Shu Fang, and Xiaowu Chen</i>	
A Matrix Decomposition Perspective to Multiple Graph Matching	199
<i>Junchi Yan, Hongteng Xu, Hongyuan Zha, Xiaokang Yang, Huanxi Liu, and Stephen Chu</i>	
Fast and Effective L0 Gradient Minimization by Region Fusion	208
<i>Rang M. H. Nguyen and Michael S. Brown</i>	
Generic Promotion of Diffusion-Based Salient Object Detection	217
<i>Peng Jiang, Nuno Vasconcelos, and Jingliang Peng</i>	

Nighttime Haze Removal with Glow and Multiple Light Colors	226
<i>Yu Li, Robby T. Tan, and Michael S. Brown</i>	
Conformal and Low-Rank Sparse Representation for Image Restoration	235
<i>Jianwei Li, Xiaowu Chen, Dongqing Zou, Bo Gao, and Wei Teng</i>	
Patch Group Based Nonlocal Self-Similarity Prior Learning for Image Denoising	244
<i>Jun Xu, Lei Zhang, Wangmeng Zuo, David Zhang, and Xiangchu Feng</i>	
Automatic Thumbnail Generation Based on Visual Representativeness and Foreground Recognizability	253
<i>Jingwei Huang, Huarong Chen, Bin Wang, and Stephen Lin</i>	
SALICON: Reducing the Semantic Gap in Saliency Prediction by Adapting Deep Neural Networks	262
<i>Xun Huang, Chengyao Shen, Xavier Boix, and Qi Zhao</i>	
A Novel Sparsity Measure for Tensor Recovery	271
<i>Qian Zhao, Deyu Meng, Xu Kong, Qi Xie, Wenfei Cao, Yao Wang, and Zongben Xu</i>	
Oriented Object Proposals	280
<i>Shengfeng He and Rynson W. H. Lau</i>	
Learning Nonlinear Spectral Filters for Color Image Reconstruction	289
<i>Michael Moeller, Julia Diebold, Guy Gilboa, and Daniel Cremers</i>	
Beyond White: Ground Truth Colors for Color Constancy Correction	298
<i>Dongliang Cheng, Brian Price, Scott Cohen, and Michael S. Brown</i>	
RGB-Guided Hyperspectral Image Upsampling	307
<i>Hyeokhyen Kwon and Yu-Wing Tai</i>	
Projection onto the Manifold of Elongated Structures for Accurate Extraction	316
<i>Amos Sironi, Vincent Lepetit, and Pascal Fua</i>	
Naive Bayes Super-Resolution Forest	325
<i>Jordi Salvador and Eduardo Pérez-Pellitero</i>	
POP Image Fusion — Derivative Domain Image Fusion without Reintegration	334
<i>Graham D. Finlayson and Alex E. Hayes</i>	
Adaptive Spatial-Spectral Dictionary Learning for Hyperspectral Image Denoising	343
<i>Ying Fu, Antony Lam, Imari Sato, and Yoichi Sato</i>	
Fully Connected Guided Image Filtering	352
<i>Longquan Dai, Mengke Yuan, Feihu Zhang, and Xiaopeng Zhang</i>	
Segment Graph Based Image Filtering: Fast Structure-Preserving Smoothing	361
<i>Feihu Zhang, Longquan Dai, Shiming Xiang, and Xiaopeng Zhang</i>	
Deep Networks for Image Super-Resolution with Sparse Prior	370
<i>Zhaowen Wang, Ding Liu, Jianchao Yang, Wei Han, and Thomas Huang</i>	
Convolutional Color Constancy	379
<i>Jonathan T. Barron</i>	
Learning Ordinal Relationships for Mid-Level Vision	388
<i>Daniel Zoran, Phillip Isola, Dilip Krishnan, and William T. Freeman</i>	

Thin Structure Estimation with Curvature Regularization	397
<i>Dmitrii Marin, Yuchen Zhong, Maria Drangova, and Yuri Boykov</i>	
HARF: Hierarchy-Associated Rich Features for Salient Object Detection	406
<i>Wenbin Zou and Nikos Komodakis</i>	
Deep Colorization	415
<i>Zezhou Cheng, Qingxiong Yang, and Bin Sheng</i>	
Image Matting with KL-Divergence Based Sparse Sampling	424
<i>Levent Karacan, Aykut Erdem, and Erkut Erdem</i>	
Intrinsic Decomposition of Image Sequences from Local Temporal Variations	433
<i>Pierre-Yves Laffont and Jean-Charles Bazin</i>	
Low-Rank Tensor Approximation with Laplacian Scale Mixture Modeling for Multiframe	
Image Denoising	442
<i>Weisheng Dong, Guangyu Li, Guangming Shi, Xin Li, and Yi Ma</i>	
Learning Parametric Distributions for Image Super-Resolution: Where Patch Matching	
Meets Sparse Coding	450
<i>Yongbo Li, Weisheng Dong, Guangming Shi, and Xuemei Xie</i>	
Improving Image Restoration with Soft-Rounding	459
<i>Xing Mei, Honggang Qi, Bao-Gang Hu, and Siwei Lyu</i>	
See the Difference: Direct Pre-Image Reconstruction and Pose Estimation	
by Differentiating HOG	468
<i>Wei-Chen Chiu and Mario Fritz</i>	
An Efficient Statistical Method for Image Noise Level Estimation	477
<i>Guangyong Chen, Fengyuan Zhu, and Pheng Ann Heng</i>	
Contour Detection and Characterization for Asynchronous Event Sensors	486
<i>Francisco Barranco, Ching L. Teo, Cornelia Fermüller, and Yiannis Aloimonos</i>	
Class-Specific Image Deblurring	495
<i>Saeed Anwar, Cong Phuoc Huynh, and Fatih Porikli</i>	
High-for-Low and Low-for-High: Efficient Boundary Detection from Deep Object	
Features and Its Applications to High-Level Vision	504
<i>Gedas Bertasius, Jianbo Shi, and Lorenzo Torresani</i>	
Variational Depth Superresolution Using Example-Based Edge Representations	513
<i>David Ferstl, Matthias R��ther, and Horst Bischof</i>	
Conditioned Regression Models for Non-blind Single Image Super-Resolution	522
<i>Gernot Riegler, Samuel Schuster, Matthias R��ther, and Horst Bischof</i>	
Video Super-Resolution via Deep Draft-Ensemble Learning	531
<i>Renjie Liao, Xin Tao, Ruiyu Li, Ziyang Ma, and Jiaya Jia</i>	
Pan-Sharpening with a Hyper-Laplacian Penalty	540
<i>Yiyong Jiang, Xinghao Ding, Delu Zeng, Yue Huang, and John Paisley</i>	
Video Restoration Against Yin-Yang Phasing	549
<i>Xiaolin Wu, Zhenhao Li, and Xiaowei Deng</i>	

Rolling Shutter Super-Resolution	558
<i>Abhijith Punnappurath, Vijay Rengarajan, and A. N. Rajagopalan</i>	
Learning Large-Scale Automatic Image Colorization	567
<i>Aditya Deshpande, Jason Rock, and David Forsyth</i>	
Compression Artifacts Reduction by a Deep Convolutional Network	576
<i>Chao Dong, Yubin Deng, Chen Change Loy, and Xiaoou Tang</i>	
Multiple-Hypothesis Affine Region Estimation with Anisotropic LoG Filters	585
<i>Takahiro Hasegawa, Mitsuru Ambai, Kohta Ishikawa, Gou Koutaki, Yuji Yamauchi, Takayoshi Yamashita, and Hironobu Fujiyoshi</i>	
A Self-Paced Multiple-Instance Learning Framework for Co-Saliency Detection	594
<i>Dingwen Zhang, Deyu Meng, Chao Li, Lu Jiang, Qian Zhao, and Junwei Han</i>	
External Patch Prior Guided Internal Clustering for Image Denoising	603
<i>Fei Chen, Lei Zhang, and Huimin Yu</i>	
Self-Calibration of Optical Lenses	612
<i>Michael Hirsch and Bernhard Schölkopf</i>	
Illumination Robust Color Naming via Label Propagation	621
<i>Yuanliu Liu, Zejian Yuan, Badong Chen, Jianru Xue, and Nanning Zheng</i>	
Unsupervised Cross-Modal Synthesis of Subject-Specific Scans	630
<i>Raviteja Vemulapalli, Hien Van Nguyen, and Shaohua Kevin Zhou</i>	
Learning to Boost Filamentary Structure Segmentation	639
<i>Lin Gu and Li Cheng</i>	
Weakly-Supervised Structured Output Learning with Flexible and Latent Graphs Using High-Order Loss Functions	648
<i>Gustavo Carneiro, Tingying Peng, Christine Bayer, and Nassir Navab</i>	
Efficient Classifier Training to Minimize False Merges in Electron Microscopy Segmentation	657
<i>Toufiq Parag, Dan C. Ciresan, and Alessandro Giusti</i>	
On Statistical Analysis of Neuroimages with Imperfect Registration	666
<i>Won Hwa Kim, Sathya N. Ravi, Sterling C. Johnson, Ozioma C. Okonkwo, and Vikas Singh</i>	

S-1A: Plenary Session

Convex Optimization with Abstract Linear Operators	675
<i>Steven Diamond and Stephen Boyd</i>	

P-1B: Recognition and 3D Computer Vision I

Building Dynamic Cloud Maps from the Ground Up	684
<i>Calvin Murdock, Nathan Jacobs, and Robert Pless</i>	
A Versatile Learning-Based 3D Temporal Tracker: Scalable, Robust, Online	693
<i>David Joseph Tan, Federico Tombari, Slobodan Ilic, and Nassir Navab</i>	
Realtime Edge-Based Visual Odometry for a Monocular Camera	702
<i>Juan José Tarrio and Sol Pedre</i>	

Fill and Transfer: A Simple Physics-Based Approach for Containability Reasoning	711
<i>Lap-Fai Yu, Noah Duncan, and Sai-Kit Yeung</i>	
On Linear Structure from Motion for Light Field Cameras	720
<i>Ole Johannsen, Antonin Sulc, and Bastian Goldluecke</i>	
3D Object Reconstruction from Hand-Object Interactions	729
<i>Dimitrios Tzionas and Juergen Gall</i>	
Minimal Solvers for 3D Geometry from Satellite Imagery	738
<i>Enliang Zheng, Ke Wang, Enrique Dunn, and Jan-Michael Frahm</i>	
An Efficient Minimal Solution for Multi-camera Motion	747
<i>Jonathan Ventura, Clemens Arth, and Vincent Lepetit</i>	
Learning Shape, Motion and Elastic Models in Force Space	756
<i>Antonio Agudo and Francesc Moreno-Noguer</i>	
A Versatile Scene Model with Differentiable Visibility Applied to Generative Pose Estimation	765
<i>Helge Rhodin, Nadia Robertini, Christian Richardt, Hans-Peter Seidel, and Christian Theobalt</i>	
Semantic Pose Using Deep Networks Trained on Synthetic RGB-D	774
<i>Jeremie Papon and Markus Schoeler</i>	
Exploiting High Level Scene Cues in Stereo Reconstruction	783
<i>Simon Hadfield and Richard Bowden</i>	
Point Triangulation through Polyhedron Collapse Using the ℓ_1 Norm	792
<i>Simon Donné, Bart Goossens, and Wilfried Philips</i>	
Optimizing the Viewing Graph for Structure-from-Motion	801
<i>Chris Sweeney, Torsten Sattler, Tobias Höllerer, Matthew Turk, and Marc Pollefeys</i>	
Intrinsic Scene Decomposition from RGB-D Images	810
<i>Mohammed Hachama, Bernard Ghanem, and Peter Wonka</i>	
3D Hand Pose Estimation Using Randomized Decision Forest with Segmentation Index Points	819
<i>Peiyi Li, Haibin Ling, Xi Li, and Chunyuan Liao</i>	
Accurate Camera Calibration Robust to Defocus Using a Smartphone	828
<i>Hyowon Ha, Yunsu Bok, Kyungdon Joo, Jiyoung Jung, and In So Kweon</i>	
High Quality Structure from Small Motion for Rolling Shutter Cameras	837
<i>Sunghoon Im, Hyowon Ha, Gyeongmin Choe, Hae-Gon Jeon, Kyungdon Joo, and In So Kweon</i>	
Photogeometric Scene Flow for High-Detail Dynamic 3D Reconstruction	846
<i>Paulo F. U. Gotardo, Tomas Simon, Yaser Sheikh, and Iain Matthews</i>	
Blur-Aware Disparity Estimation from Defocus Stereo Images	855
<i>Ching-Hui Chen, Hui Zhou, and Timo Ahonen</i>	
Global Structure-from-Motion by Similarity Averaging	864
<i>Zhaopeng Cui and Ping Tan</i>	

Massively Parallel Multiview Stereopsis by Surface Normal Diffusion	873
<i>Silvano Galliani, Katrin Lasinger, and Konrad Schindler</i>	
Variational PatchMatch MultiView Reconstruction and Refinement	882
<i>Philipp Heise, Brian Jensen, Sebastian Klose, and Alois Knoll</i>	
As-Rigid-as-Possible Volumetric Shape-from-Template	891
<i>Shaifali Parashar, Daniel Pizarro, Adrien Bartoli, and Toby Collins</i>	
General Dynamic Scene Reconstruction from Multiple View Video	900
<i>Armin Mustafa, Hansung Kim, Jean-Yves Guillemaut, and Adrian Hilton</i>	
The Joint Image Handbook	909
<i>Matthew Trager, Martial Hebert, and Jean Ponce</i>	
Direct, Dense, and Deformable: Template-Based Non-rigid 3D Reconstruction from RGB Video	918
<i>Rui Yu, Chris Russell, Neill D. F. Campbell, and Lourdes Agapito</i>	
Single Image Pop-Up from Discriminatively Learned Parts	927
<i>Menglong Zhu, Xiaowei Zhou, and Kostas Daniilidis</i>	
Learning Informative Edge Maps for Indoor Scene Layout Prediction	936
<i>Arun Mallya and Svetlana Lazebnik</i>	
Multi-view Convolutional Neural Networks for 3D Shape Recognition	945
<i>Hang Su, Subhransu Maji, Evangelos Kalogerakis, and Erik Learned-Miller</i>	
Learning Analysis-by-Synthesis for 6D Pose Estimation in RGB-D Images	954
<i>Alexander Krull, Eric Brachmann, Frank Michel, Michael Ying Yang, Stefan Gumhold, and Carsten Rother</i>	
3D Surface Profilometry Using Phase Shifting of De Bruijn Pattern	963
<i>Matea Đonlić, Tomislav Petković, and Tomislav Pribanić</i>	
A Deep Visual Correspondence Embedding Model for Stereo Matching Costs	972
<i>Zhuoyuan Chen, Xun Sun, Liang Wang, Yinan Yu, and Chang Huang</i>	
Learning Concept Embeddings with Combined Human-Machine Expertise	981
<i>Michael J. Wilber, Iljung S. Kwak, David Kriegman, and Serge Belongie</i>	
Deep Multi-patch Aggregation Network for Image Style, Aesthetics, and Quality Estimation	990
<i>Xin Lu, Zhe Lin, Xiaohui Shen, Radomír Měch, and James Z. Wang</i>	
Towards Computational Baby Learning: A Weakly-Supervised Approach for Object Detection	999
<i>Xiaodan Liang, Si Liu, Yunchao Wei, Luoqi Liu, Liang Lin, and Shuicheng Yan</i>	
Improving Image Classification with Location Context	1008
<i>Kevin Tang, Manohar Paluri, Li Fei-Fei, Rob Fergus, and Lubomir Bourdev</i>	
HICO: A Benchmark for Recognizing Human-Object Interactions in Images	1017
<i>Yu-Wei Chao, Zhan Wang, Yugeng He, Jiaxuan Wang, and Jia Deng</i>	
Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification	1026
<i>Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun</i>	

Continuous Pose Estimation with a Spatial Ensemble of Fisher Regressors	1035
<i>Michele Fenzi, Laura Leal-Taixé, Jörn Ostermann, and Tinne Tuytelaars</i>	
Adaptive Hashing for Fast Similarity Search	1044
<i>Fatih Cakir and Stan Sclaroff</i>	
Single Image 3D without a Single 3D Image	1053
<i>David F. Fouhey, Wajahat Hussain, Abhinav Gupta, and Martial Hebert</i>	
Cross-Domain Image Retrieval with a Dual Attribute-Aware Ranking Network	1062
<i>Junshi Huang, Rogerio Feris, Qiang Chen, and Shuicheng Yan</i>	
Attribute-Graph: A Graph Based Approach to Image Ranking	1071
<i>Nikita Prabhu and R. Venkatesh Babu</i>	
Contextual Action Recognition with R [*] CNN	1080
<i>Georgia Gkioxari, Ross Girshick, and Jitendra Malik</i>	
What Makes an Object Memorable?	1089
<i>Rachit Dubey, Joshua Peterson, Aditya Khosla, Ming-Hsuan Yang, and Bernard Ghanem</i>	
kNN Hashing with Factorized Neighborhood Representation	1098
<i>Kun Ding, Chunlei Huo, Bin Fan, and Chunhong Pan</i>	
Multi-View Complementary Hash Tables for Nearest Neighbor Search	1107
<i>Xianglong Liu, Lei Huang, Cheng Deng, Jiwen Lu, and Bo Lang</i>	
Scalable Person Re-identification: A Benchmark	1116
<i>Liang Zheng, Liyue Shen, Lu Tian, Shengjin Wang, Jingdong Wang, and Qi Tian</i>	
MMSS: Multi-modal Sharable and Specific Feature Learning for RGB-D Object Recognition	1125
<i>Anran Wang, Jianfei Cai, Jiwen Lu, and Tat-Jen Cham</i>	
Object Detection via a Multi-region and Semantic Segmentation-Aware CNN Model	1134
<i>Spyros Gidaris and Nikos Komodakis</i>	
Neural Activation Constellations: Unsupervised Part Model Discovery with Convolutional Networks	1143
<i>Marcel Simon and Erik Rodner</i>	
Cascaded Sparse Spatial Bins for Efficient and Effective Generic Object Detection	1152
<i>David Novotny and Jiri Matas</i>	
Probabilistic Label Relation Graphs with Ising Models	1161
<i>Nan Ding, Jia Deng, Kevin P. Murphy, and Hartmut Neven</i>	
Predicting Good Features for Image Geo-Localization Using Per-Bundle VLAD	1170
<i>Hyo Jin Kim, Enrique Dunn, and Jan-Michael Frahm</i>	
Task-Driven Feature Pooling for Image Classification	1179
<i>Guo-Sen Xie, Xu-Yao Zhang, Xiangbo Shu, Shuicheng Yan, and Cheng-Lin Liu</i>	
Cutting Edge: Soft Correspondences in Multimodal Scene Parsing	1188
<i>Sarah Taghavi Namin, Mohammad Najafi, Mathieu Salzmann, and Lars Petersson</i>	
One Shot Learning via Compositions of Meaningful Patches	1197
<i>Alex Wong and Alan Yuille</i>	

FASText: Efficient Unconstrained Scene Text Detector	1206
<i>Michal Buřta, Lukáš Neumann, and Jiří Matas</i>	
Multi-scale Recognition with DAG-CNNs	1215
<i>Songfan Yang and Deva Ramanan</i>	
Relaxed Multiple-Instance SVM with Application to Object Discovery	1224
<i>Xinggang Wang, Zhuotun Zhu, Cong Yao, and Xiang Bai</i>	
Im2Calories: Towards an Automated Mobile Vision Food Diary	1233
<i>Austin Myers, Nick Johnston, Vivek Rathod, Anoop Korattikara, Alex Gorban, Nathan Silberman, Sergio Guadarrama, George Papandreou, Jonathan Huang, and Kevin Murphy</i>	
LEWIS: Latent Embeddings for Word Images and Their Semantics	1242
<i>Albert Gordo, Jon Almazán, Naila Murray, and Florent Perronin</i>	
Per-Sample Kernel Adaptation for Visual Recognition and Grouping	1251
<i>Borislav Antic and Björn Ommer</i>	
Fine-Grained Change Detection of Misaligned Scenes with Varied Illuminations	1260
<i>Wei Feng, Fei-Peng Tian, Qian Zhang, Nan Zhang, Liang Wan, and Jizhou Sun</i>	
Aggregating Local Deep Features for Image Retrieval	1269
<i>Artem Babenko Yandex and Victor Lempitsky</i>	
Learning Deep Object Detectors from 3D Models	1278
<i>Xingchao Peng, Baochen Sun, Karim Ali, and Kate Saenko</i>	
Harvesting Discriminative Meta Objects with Deep CNN Features for Scene Classification	1287
<i>Ruobing Wu, Baoyuan Wang, Wenping Wang, and Yizhou Yu</i>	
Scalable Nonlinear Embeddings for Semantic Category-Based Image Retrieval	1296
<i>Gaurav Sharma and Bernt Schiele</i>	
Person Re-Identification Ranking Optimisation by Discriminant Context Information Analysis	1305
<i>Jorge García, Niki Martinel, Christian Micheloni, and Alfredo Gardel</i>	
Unsupervised Generation of a View Point Annotated Car Dataset from Videos	1314
<i>Nima Sedaghat and Thomas Brox</i>	

O-1B: 3D Vision

Structured Indoor Modeling	1323
<i>Satoshi Ikehata, Hang Yang, and Yasutaka Furukawa</i>	
3D Time-Lapse Reconstruction from Internet Photos	1332
<i>Ricardo Martin-Brualla, David Gallup, and Steven M. Seitz</i>	
Global, Dense Multiscale Reconstruction for a Billion Points	1341
<i>Benjamin Ummerhofer and Thomas Brox</i>	
On the Visibility of Point Clouds	1350
<i>Sagi Katz and Ayellet Tal</i>	

O-2A: Segmentation, Edges and Saliency

Weakly Supervised Graph Based Semantic Segmentation by Learning Communities of Image-Parts	1359
<i>Niloufar Pourian, S. Karthikeyan, and B. S. Manjunath</i>	
Piecewise Flat Embedding for Image Segmentation	1368
<i>Yizhou Yu, Chaowei Fang, and Zicheng Liao</i>	
Semantic Image Segmentation via Deep Parsing Network	1377
<i>Ziwei Liu, Xiaoxiao Li, Ping Luo, Chen-Change Loy, and Xiaoou Tang</i>	
Human Parsing with Contextualized Convolutional Neural Network	1386
<i>Xiaodan Liang, Chunyan Xu, Xiaohui Shen, Jianchao Yang, Si Liu, Jinhui Tang, Liang Lin, and Shuicheng Yan</i>	
Holistically-Nested Edge Detection	1395
<i>Saining Xie and Zhuowen Tu</i>	
Minimum Barrier Salient Object Detection at 80 FPS	1404
<i>Jianming Zhang, Stan Sclaroff, Zhe Lin, Xiaohui Shen, Brian Price, and Radomír Měch</i>	

O-2B: Learning Representations and Attributes

Learning Image Representations Tied to Ego-Motion	1413
<i>Dinesh Jayaraman and Kristen Grauman</i>	
Unsupervised Visual Representation Learning by Context Prediction	1422
<i>Carl Doersch, Abhinav Gupta, and Alexei A. Efros</i>	
Webly Supervised Learning of Convolutional Networks	1431
<i>Xinlei Chen and Abhinav Gupta</i>	
Fast R-CNN	1440
<i>Ross Girshick</i>	
Bilinear CNN Models for Fine-Grained Visual Recognition	1449
<i>Tsung-Yu Lin, Aruni RoyChowdhury, and Subhansu Maji</i>	
Discovering the Spatial Extent of Relative Attributes	1458
<i>Fanyi Xiao and Yong Jae Lee</i>	

O-2C: Statistical Methods and Learning

Deep Neural Decision Forests	1467
<i>Peter Kotschieder, Madalina Fiterau, Antonio Criminisi, and Samuel Rota Bulò</i>	
Deep Fried Convnets	1476
<i>Zichao Yang, Marcin Moczulski, Misha Denil, Nando de Freitas, Alex Smola, Le Song, and Ziyu Wang</i>	
Semantic Component Analysis	1484
<i>Calvin Murdock and Fernando De la Torre</i>	

Low-Rank Matrix Factorization under General Mixture Noise Distributions	1493
<i>Xiangyong Cao, Yang Chen, Qian Zhao, Deyu Meng, Yao Wang, Dong Wang, and Zongben Xu</i>	
Web-Scale Image Clustering Revisited	1502
<i>Yannis Avrithis, Yannis Kalantidis, Evangelos Anagnostopoulos, and Ioannis Z. Emiris</i>	
Learning Discriminative Reconstructions for Unsupervised Outlier Removal	1511
<i>Yan Xia, Xudong Cao, Fang Wen, Gang Hua, and Jian Sun</i>	
P-2A: Optimization, Segmentation, and Recognition	
Learning Deconvolution Network for Semantic Segmentation	1520
<i>Hyeonwoo Noh, Seunghoon Hong, and Bohyung Han</i>	
Conditional Random Fields as Recurrent Neural Networks	1529
<i>Shuai Zheng, Sadeep Jayasumana, Bernardino Romera-Paredes, Vibhav Vineet, Zhizhong Su, Dalong Du, Chang Huang, and Philip H. S. Torr</i>	
The One Triangle Three Parallelograms Sampling Strategy and Its Application in Shape Regression	1538
<i>Mikael Nilsson</i>	
Boosting Object Proposals: From Pascal to COCO	1546
<i>Jordi Pont-Tuset and Luc Van Gool</i>	
Secrets of GrabCut and Kernel K-Means	1555
<i>Meng Tang, Ismail Ben Ayed, Dmitrii Marin, and Yuri Boykov</i>	
Video Matting via Sparse and Low-Rank Representation	1564
<i>Dongqing Zou, Xiaowu Chen, Guangying Cao, and Xiaogang Wang</i>	
Joint Object and Part Segmentation Using Deep Learned Potentials	1573
<i>Peng Wang, Xiaohui Shen, Zhe Lin, Scott Cohen, Brian Price, and Alan Yuille</i>	
Low-Rank Tensor Constrained Multiview Subspace Clustering	1582
<i>Changqing Zhang, Huazhu Fu, Si Liu, Guangcan Liu, and Xiaochun Cao</i>	
BodyPrint: Pose Invariant 3D Shape Matching of Human Bodies	1591
<i>Jiangping Wang, Kai Ma, Vivek Kumar Singh, Thomas Huang, and Terrence Chen</i>	
The Middle Child Problem: Revisiting Parametric Min-Cut and Seeds for Object Proposals	1600
<i>Ahmad Humayun, Fuxin Li, and James M. Rehg</i>	
Contour Guided Hierarchical Model for Shape Matching	1609
<i>Yuanqi Su, Yuehu Liu, Bonan Cuan, and Nanning Zheng</i>	
Robust Image Segmentation Using Contour-Guided Color Palettes	1618
<i>Xiang Fu, Chien-Yi Wang, Chen Chen, Changhu Wang, and C.-C. Jay Kuo</i>	
Joint Optimization of Segmentation and Color Clustering	1626
<i>Ekaterina Lobacheva, Olga Veksler, and Yuri Boykov</i>	
BoxSup: Exploiting Bounding Boxes to Supervise Convolutional Networks for Semantic Segmentation	1635
<i>Jifeng Dai, Kaiming He, and Jian Sun</i>	

Detection and Segmentation of 2D Curved Reflection Symmetric Structures	1644
<i>Ching L. Teo, Cornelia Fermüller, and Yiannis Aloimonos</i>	
Unsupervised Tube Extraction Using Transductive Learning and Dense Trajectories	1653
<i>Mihai Marian Puscas, Enver Sangineto, Dubravko Culibrk, and Nicu Sebe</i>	
Compositional Hierarchical Representation of Shape Manifolds for Classification of Non-manifold Shapes	1662
<i>Mete Ozay, Umit Rusen Aktas, Jeremy L. Wyatt, and Aleš Leonardis</i>	
Shell PCA: Statistical Shape Modelling in Shell Space	1671
<i>Chao Zhang, Behrend Heeren, Martin Rumpf, and William A. P. Smith</i>	
Learning to Combine Mid-Level Cues for Object Proposal Generation	1680
<i>Tom Lee, Sanja Fidler, and Sven Dickinson</i>	
Enhancing Road Maps by Parsing Aerial Images Around the World	1689
<i>Gellért Mátyus, Shenlong Wang, Sanja Fidler, and Raquel Urtasun</i>	
Probabilistic Appearance Models for Segmentation and Classification	1698
<i>Julia Krüger, Jan Ehrhardt, and Heinz Handels</i>	
A Randomized Ensemble Approach to Industrial CT Segmentation	1707
<i>Hyojin Kim, Jayaraman J. Thiagarajan, and Peer-Timo Bremer</i>	
Semi-Supervised Normalized Cuts for Image Segmentation	1716
<i>Selene E. Chew and Nathan D. Cahill</i>	
StereoSnakes: Contour Based Consistent Object Extraction for Stereo Images	1724
<i>Ran Ju, Tongwei Ren, and Gangshan Wu</i>	
Semantic Segmentation of RGBD Images with Mutex Constraints	1733
<i>Zhuo Deng, Sinisa Todorovic, and Longin Jan Latecki</i>	
Weakly-and Semi-Supervised Learning of a Deep Convolutional Network for Semantic Image Segmentation	1742
<i>George Papandreou, Liang-Chieh Chen, Kevin P. Murphy, and Alan L. Yuille</i>	
Efficient Decomposition of Image and Mesh Graphs by Lifted Multicuts	1751
<i>M. Keuper, E. Levinkov, N. Bonneel, G. Lavoué, T. Brox, and B. Andres</i>	
Parsimonious Labeling	1760
<i>Puneet K. Dokania and M. Pawan Kumar</i>	
Volumetric Bias in Segmentation and Reconstruction: Secrets and Solutions	1769
<i>Yuri Boykov, Hossam Isack, Carl Olsson, and Ismail Ben Ayed</i>	
Entropy Minimization for Convex Relaxation Approaches	1778
<i>Mohamed Souiai, Martin R. Oswald, Youngwook Kee, Junmo Kim, Marc Pollefeys, and Daniel Cremers</i>	
Adaptively Unified Semi-Supervised Dictionary Learning with Active Points	1787
<i>Xiaobo Wang, Xiaojie Guo, and Stan Z. Li</i>	
Constrained Convolutional Neural Networks for Weakly Supervised Segmentation	1796
<i>Deepak Pathak, Philipp Krähenbühl, and Trevor Darrell</i>	
A Multiscale Variable-Grouping Framework for MRF Energy Minimization	1805
<i>Omer Meir, Meirav Galun, Stav Yagev, Ronen Basri, and Irad Yavneh</i>	

Inferring M-Best Diverse Labelings in a Single One	1814
<i>Alexander Kirillov, Bogdan Savchynskyy, Dmitriy Schlesinger, Dmitry Vetrov, and Carsten Rother</i>	
Convolutional Sparse Coding for Image Super-Resolution	1823
<i>Shuhang Gu, Wangmeng Zuo, Qi Xie, Deyu Meng, Xiangchu Feng, and Lei Zhang</i>	
A Wavefront Marching Method for Solving the Eikonal Equation on Cartesian Grids	1832
<i>B. Cancela, M. Ortega, and M. G. Penedo</i>	
A Projection Free Method for Generalized Eigenvalue Problem with a Nonsmooth Regularizer	1841
<i>Seong Jae Hwang, Maxwell D. Collins, Sathya N. Ravi, Vamsi K. Ithapu, Nagesh Adluru, Sterling C. Johnson, and Vikas Singh</i>	
Optimizing Expected Intersection-Over-Union with Candidate-Constrained CRFs	1850
<i>Faruk Ahmed, Dany Tarlow, and Dhruv Batra</i>	
Higher-Order Inference for Multi-class Log-Supermodular Models	1859
<i>Jian Zhang, Josip Djolonga, and Andreas Krause</i>	
Depth-Based Hand Pose Estimation: Data, Methods, and Challenges	1868
<i>James S. Supančič III, Grégory Rogez, Yi Yang, Jamie Shotton, and Deva Ramanan</i>	
Adaptive Dither Voting for Robust Spatial Verification	1877
<i>Xiaomeng Wu and Kunio Kashino</i>	
Alternating Co-Quantization for Cross-Modal Hashing	1886
<i>Go Irie, Hiroyuki Arai, and Yukinobu Taniguchi</i>	
Learning Deep Representation with Large-Scale Attributes	1895
<i>Wanli Ouyang, Hongyang Li, Xingyu Zeng, and Xiaogang Wang</i>	
Deep Learning Strong Parts for Pedestrian Detection	1904
<i>Yonglong Tian, Ping Luo, Xiaogang Wang, and Xiaoou Tang</i>	
Flowing ConvNets for Human Pose Estimation in Videos	1913
<i>Tomas Pfister, James Charles, and Andrew Zisserman</i>	
Top Rank Supervised Binary Coding for Visual Search	1922
<i>Dongjin Song, Wei Liu, Rongrong Ji, David A. Meyer, and John R. Smith</i>	
BubbLeNet: Foveated Imaging for Visual Discovery	1931
<i>Kevin Matzen and Noah Snavely</i>	
PQTable: Fast Exact Asymmetric Distance Neighbor Search for Product Quantization Using Hash Tables	1940
<i>Yusuke Matsui, Toshihiko Yamasaki, and Kiyoharu Aizawa</i>	
Lending A Hand: Detecting Hands and Recognizing Activities in Complex Egocentric Interactions	1949
<i>Sven Bambach, Stefan Lee, David J. Crandall, and Chen Yu</i>	
Fast and Accurate Head Pose Estimation via Random Projection Forests	1958
<i>Donghoon Lee, Ming-Hsuan Yang, and Songhwai Oh</i>	
An MRF-Poselets Model for Detecting Highly Articulated Humans	1967
<i>Duc Thanh Nguyen, Minh-Khoi Tran, and Sai-Kit Yeung</i>	

Beyond Tree Structure Models: A New Occlusion Aware Graphical Model for Human Pose Estimation	1976
<i>Lianrui Fu, Junge Zhang, and Kaiqi Huang</i>	
Relaxing from Vocabulary: Robust Weakly-Supervised Deep Learning for Vocabulary-Free Image Tagging	1985
<i>Jianlong Fu, Yue Wu, Tao Mei, Jinqiao Wang, Hanqing Lu, and Yong Rui</i>	
Visual Phrases for Exemplar Face Detection	1994
<i>Vijay Kumar, Anoop Namboodiri, and C. V. Jawahar</i>	
Spatial Semantic Regularisation for Large Scale Object Detection	2003
<i>Damian Mrowca, Marcus Rohrbach, Judy Hoffman, Ronghang Hu, Kate Saenko, and Trevor Darrell</i>	
Human Pose Estimation in Videos	2012
<i>Dong Zhang and Mubarak Shah</i>	
Contour Box: Rejecting Object Proposals without Explicit Closed Contours	2021
<i>Cewu Lu, Shu Liu, Jiaya Jia, and Chi-Keung Tang</i>	

O-3A: Registration, Alignment and Stereo

Registering Images to Untextured Geometry Using Average Shading Gradients	2030
<i>Tobias Plötz and Stefan Roth</i>	
Robust Nonrigid Registration by Convex Optimization	2039
<i>Qifeng Chen and Vladlen Koltun</i>	
Robust and Optimal Sum-of-Squares-Based Point-to-Plane Registration of Image Sets and Structured Scenes	2048
<i>Danda Pani Paudel, Adlane Habed, Cédric Demonceaux, and Pascal Vasseur</i>	
MeshStereo: A Global Stereo Model with Mesh Alignment Regularization for View Interpolation	2057
<i>Chi Zhang, Zhiwei Li, Yanhua Cheng, Rui Cai, Hongyang Chao, and Yong Rui</i>	
CV-HAZOP: Introducing Test Data Validation for Computer Vision	2066
<i>Oliver Zendel, Markus Murschitz, Martin Humenberger, and Wolfgang Herzner</i>	

P-3A: Recognition and 3D Computer Vision II

Structure from Motion Using Structure-Less Resection	2075
<i>Enliang Zheng and Changchang Wu</i>	
Joint Camera Clustering and Surface Segmentation for Large-Scale Multi-view Stereo	2084
<i>Runze Zhang, Shiwei Li, Tian Fang, Siyu Zhu, and Long Quan</i>	
Higher-Order CRF Structural Segmentation of 3D Reconstructed Surfaces	2093
<i>Jingbo Liu, Jinglu Wang, Tian Fang, Chiew-Lan Tai, and Long Quan</i>	
Hyperpoints and Fine Vocabularies for Large-Scale Location Recognition	2102
<i>Torsten Sattler, Michal Havlena, Filip Radenović, Konrad Schindler, and Marc Pollefeys</i>	
Globally Optimal 2D-3D Registration from Points or Lines without Correspondences	2111
<i>Mark Brown, David Windridge, and Jean-Yves Guillemaut</i>	

The HCI Stereo Metrics: Geometry-Aware Performance Analysis of Stereo Algorithms	2120
<i>Katrin Honauer, Lena Maier-Hein, and Daniel Kondermann</i>	
Merging the Unmatchable: Stitching Visually Disconnected SfM Models	2129
<i>Andrea Cohen, Torsten Sattler, and Marc Pollefeys</i>	
3D Fragment Reassembly Using Integrated Template Guidance and Fracture-Region Matching	2138
<i>Kang Zhang, Wuyi Yu, Mary Manhein, Warren Waggenspack, and Xin Li</i>	
Procedural Editing of 3D Building Point Clouds	2147
<i>Ilke Demir, Daniel G. Aliaga, and Bedrich Benes</i>	
Semantically-Aware Aerial Reconstruction from Multi-modal Data	2156
<i>Randi Cabezas, Julian Straub, and John W. Fisher III</i>	
Guaranteed Outlier Removal for Rotation Search	2165
<i>Álvaro Parra Bustos and Tat-Jun Chin</i>	
Peeking Template Matching for Depth Extension	2174
<i>Simon Korman, Eyal Ofek, and Shai Avidan</i>	
Deformable 3D Fusion: From Partial Dynamic 3D Observations to Complete 4D Models	2183
<i>Weipeng Xu, Mathieu Salzmann, Yongtian Wang, and Yue Liu</i>	
Non-parametric Structure-Based Calibration of Radially Symmetric Cameras	2192
<i>Federico Camposeco, Torsten Sattler, and Marc Pollefeys</i>	
Exploiting Object Similarity in 3D Reconstruction	2201
<i>Chen Zhou, Fatma Güney, Yizhou Wang, and Andreas Geiger</i>	
You are Here: Mimicking the Human Thinking Process in Reading Floor-Plans	2210
<i>Hang Chu, Dong Ki Kim, and Tsuhan Chen</i>	
MAP Disparity Estimation Using Hidden Markov Trees	2219
<i>Eric T. Psota, Jędrzej Kowalczyk, Mateusz Mittek, and Lance C. Pérez</i>	
Wide Baseline Stereo Matching with Convex Bounded Distortion Constraints	2228
<i>Meirav Galun, Tal Amir, Tal Hassner, Ronen Basri, and Yaron Lipman</i>	
Interactive Visual Hull Refinement for Specular and Transparent Object Surface Reconstruction	2237
<i>Xinxin Zuo, Chao Du, Sen Wang, Jiangbin Zheng, and Ruigang Yang</i>	
Hierarchical Higher-Order Regression Forest Fields: An Application to 3D Indoor Scene Labelling	2246
<i>Trung T. Pham, Ian Reid, Yasir Latif, and Stephen Gould</i>	
Classical Scaling Revisited	2255
<i>Gil Shamaï, Yonathan Aflalo, Michael Zibulevsky, and Ron Kimmel</i>	
Dense Continuous-Time Tracking and Mapping with Rolling Shutter RGB-D Cameras	2264
<i>Christian Kerl, Jörg Stückler, and Daniel Cremers</i>	
Dense Image Registration and Deformable Surface Reconstruction in Presence of Occlusions and Minimal Texture	2273
<i>Dat Tien Ngo, Sanghyuk Park, Anne Jorstad, Alberto Crivellaro, Chang D. Yoo, and Pascal Fua</i>	

The Likelihood-Ratio Test and Efficient Robust Estimation	2282
<i>Andrea Cohen and Christopher Zach</i>	
Reflection Modeling for Passive Stereo	2291
<i>Rahul Nair, Andrew Fitzgibbon, Daniel Kondermann, and Carsten Rother</i>	
Detailed Full-Body Reconstructions of Moving People from Monocular RGB-D Sequences	2300
<i>Federica Bogo, Michael J. Black, Matthew Loper, and Javier Romero</i>	
Efficient Solution to the Epipolar Geometry for Radially Distorted Cameras	2309
<i>Zuzana Kukelova, Jan Heller, Martin Bujnak, Andrew Fitzgibbon, and Tomas Pajdla</i>	
Learning a Descriptor-Specific 3D Keypoint Detector	2318
<i>Samuele Salti, Federico Tombari, Riccardo Spezialetti, and Luigi Di Stefano</i>	
Component-Wise Modeling of Articulated Objects	2327
<i>Valsamis Ntouskos, Marta Sanzari, Bruno Cafaro, Federico Nardi, Fabrizio Natola, Fiora Pirri, and Manuel Ruiz</i>	
A Collaborative Filtering Approach to Real-Time Hand Pose Estimation	2336
<i>Chiho Choi, Ayan Sinha, Joon Hee Choi, Sujin Jang, and Karthik Ramani</i>	
On the Equivalence of Moving Entrance Pupil and Radial Distortion for Camera Calibration	2345
<i>Avinash Kumar and Narendra Ahuja</i>	
A Linear Generalized Camera Calibration from Three Intersecting Reference Planes	2354
<i>Mai Nishimura, Shohei Nobuhara, Takashi Matsuyama, Shinya Shimizu, and Kensaku Fujii</i>	
Towards Pointless Structure from Motion: 3D Reconstruction and Camera Parameters from General 3D Curves	2363
<i>Irina Nurutdinova and Andrew Fitzgibbon</i>	
Attributed Grammars for Joint Estimation of Human Attributes, Part and Pose	2372
<i>Seyoung Park and Song-Chun Zhu</i>	
Real-Time Pose Estimation Piggybacked on Object Detection	2381
<i>Roman Juránek, Adam Herout, Markéta Dubská, and Pavel Zemčík</i>	
Understanding and Predicting Image Memorability at a Large Scale	2390
<i>Aditya Khosla, Akhil S. Raju, Antonio Torralba, and Aude Oliva</i>	
Multiple Granularity Descriptors for Fine-Grained Categorization	2399
<i>Dequan Wang, Zhiqiang Shen, Jie Shao, Wei Zhang, Xiangyang Xue, and Zheng Zhang</i>	
Guiding the Long-Short Term Memory Model for Image Caption Generation	2407
<i>Xu Jia, Efstratios Gavves, Basura Fernando, and Tinne Tuytelaars</i>	
Just Noticeable Differences in Visual Attributes	2416
<i>Aron Yu and Kristen Grauman</i>	
VQA: Visual Question Answering	2425
<i>Stanislaw Antol, Aishwarya Agrawal, Jiasen Lu, Margaret Mitchell, Dhruv Batra, C. Lawrence Zitnick, and Devi Parikh</i>	
Localize Me Anywhere, Anytime: A Multi-task Point-Retrieval Approach	2434
<i>Guoyu Lu, Yan Yan, Li Ren, Jingkuan Song, Nicu Sebe, and Chandra Kambhampettu</i>	

Dense Optical Flow Prediction from a Static Image	2443
<i>Jacob Walker, Abhinav Gupta, and Martial Hebert</i>	
Unsupervised Domain Adaptation for Zero-Shot Learning	2452
<i>Elyor Kodirov, Tao Xiang, Zhenyong Fu, and Shaogang Gong</i>	
Visual Madlibs: Fill in the Blank Description Generation and Question Answering	2461
<i>Licheng Yu, Eunbyung Park, Alexander C. Berg, and Tamara L. Berg</i>	
Actions and Attributes from Wholes and Parts	2470
<i>Georgia Gkioxari, Ross Girshick, and Jitendra Malik</i>	
DeepBox: Learning Objectness with Convolutional Networks	2479
<i>Weicheng Kuo, Bharath Hariharan, and Jitendra Malik</i>	
Active Object Localization with Deep Reinforcement Learning	2488
<i>Juan C. Caicedo and Svetlana Lazebnik</i>	
Scene-Domain Active Part Models for Object Representation	2497
<i>Zhou Ren, Chaohui Wang, and Alan Yuille</i>	
A Unified Multiplicative Framework for Attribute Learning	2506
<i>Kongming Liang, Hong Chang, Shiguang Shan, and Xilin Chen</i>	
Contractive Rectifier Networks for Nonlinear Maximum Margin Classification	2515
<i>Senjian An, Munawar Hayat, Salman H. Khan, Mohammed Bennamoun, Farid Boussaid, and Ferdous Sohel</i>	
Augmenting Strong Supervision Using Web Data for Fine-Grained Categorization	2524
<i>Zhe Xu, Shaoli Huang, Ya Zhang, and Dacheng Tao</i>	
Learning Like a Child: Fast Novel Visual Concept Learning from Sentence Descriptions of Images	2533
<i>Junhua Mao, Xu Wei, Yi Yang, Jiang Wang, Zhiheng Huang, and Alan L. Yuille</i>	
Learning Common Sense through Visual Abstraction	2542
<i>Ramakrishna Vedantam, Xiao Lin, Tanmay Batra, C. Lawrence Zitnick, and Devi Parikh</i>	
Domain Generalization for Object Recognition with Multi-task Autoencoders	2551
<i>Muhammad Ghifary, W. Bastiaan Kleijn, Mengjie Zhang, and David Balduzzi</i>	
Square Localization for Efficient and Accurate Object Detection	2560
<i>Cewu Lu, Yongyi Lu, Hao Chen, and Chi-Keung Tang</i>	
Box Aggregation for Proposal Decimation: Last Mile of Object Detection	2569
<i>Shu Liu, Cewu Lu, and Jiaya Jia</i>	
DeepProposal: Hunting Objects by Cascading Deep Convolutional Layers	2578
<i>Amir Ghodrati, Ali Diba, Marco Pedersoli, Tinne Tuytelaars, and Luc Van Gool</i>	
Semantic Segmentation with Object Clique Potential	2587
<i>Xiaojuan Qi, Jianping Shi, Shu Liu, Renjie Liao, and Jiaya Jia</i>	
Automatic Concept Discovery from Parallel Text and Visual Corpora	2596
<i>Chen Sun, Chuang Gan, and Ram Nevatia</i>	
Simpler Non-Parametric Methods Provide as Good or Better Results to Multiple-Instance Learning	2605
<i>Ragav Venkatesan, Parag Shridhar Chandakkar, and Baoxin Li</i>	

Monocular Object Instance Segmentation and Depth Ordering with CNNs	2614
<i>Ziyu Zhang, Alexander G. Schwing, Sanja Fidler, and Raquel Urtasun</i>	
Multimodal Convolutional Neural Networks for Matching Image and Sentence	2623
<i>Lin Ma, Zhengdong Lu, Lifeng Shang, and Hang Li</i>	
Structural Kernel Learning for Large Scale Multiclass Object Co-detection	2632
<i>Zeeshan Hayder, Xuming He, and Mathieu Salzmann</i>	
Flickr30k Entities: Collecting Region-to-Phrase Correspondences for Richer Image-to-Sentence Models	2641
<i>Bryan A. Plummer, Liwei Wang, Chris M. Cervantes, Juan C. Caicedo, Julia Hockenmaier, and Svetlana Lazebnik</i>	
Predicting Depth, Surface Normals and Semantic Labels with a Common Multi-scale Convolutional Architecture	2650
<i>David Eigen and Rob Fergus</i>	
AttentionNet: Aggregating Weak Directions for Accurate Object Detection	2659
<i>Donggeun Yoo, Sunggyun Park, Joon-Young Lee, Anthony S. Paek, and In So Kweon</i>	
Common Subspace for Model and Similarity: Phrase Learning for Caption Generation from Images	2668
<i>Yoshitaka Ushiku, Masataka Yamaguchi, Yusuke Mukuta, and Tatsuya Harada</i>	

O-3B: Representations for Recognition and Localization

3D-Assisted Feature Synthesis for Novel Views of an Object	2677
<i>Hao Su, Fan Wang, Eric Yi, and Leonidas Guibas</i>	
Render for CNN: Viewpoint Estimation in Images Using CNNs Trained with Rendered 3D Model Views	2686
<i>Hao Su, Charles R. Qi, Yangyan Li, and Leonidas J. Guibas</i>	
Lost Shopping! Monocular Localization in Large Indoor Spaces	2695
<i>Shenlong Wang, Sanja Fidler, and Raquel Urtasun</i>	
Camera Pose Voting for Large-Scale Image-Based Localization	2704
<i>Bernhard Zeisl, Torsten Sattler, and Marc Pollefeys</i>	

P-3B: Statistical Methods and Learning, Motion and Tracking, and Video Analysis I

MANTRA: Minimum Maximum Latent Structural SVM for Image Classification and Ranking	2713
<i>Thibaut Durand, Nicolas Thome, and Matthieu Cord</i>	
DeepDriving: Learning Affordance for Direct Perception in Autonomous Driving	2722
<i>Chenyi Chen, Ari Seff, Alain Kornhauser, and Jianxiang Xiao</i>	
Active Transfer Learning with Zero-Shot Priors: Reusing Past Datasets for Future Tasks	2731
<i>E. Gavves, T. Mensink, T. Tommasi, C. G. M. Snoek, and T. Tuytelaars</i>	

HD-CNN: Hierarchical Deep Convolutional Neural Networks for Large Scale Visual Recognition	2740
<i>Zhicheng Yan, Hao Zhang, Robinson Piramuthu, Vignesh Jagadeesh, Dennis DeCoste, Wei Di, and Yizhou Yu</i>	
Learning the Structure of Deep Convolutional Networks	2749
<i>Jiashi Feng and Trevor Darrell</i>	
FlowNet: Learning Optical Flow with Convolutional Networks	2758
<i>Alexey Dosovitskiy, Philipp Fischery, Eddy Ilg, Philip Häusser, Caner Hazırbaş, Vladimir Golkov, Patrick van der Smagt, Daniel Cremers, and Thomas Brox</i>	
Learning Semi-Supervised Representation Towards a Unified Optimization Framework for Semi-Supervised Learning	2767
<i>Chun-Guang Li, Zhouchen Lin, Honggang Zhang, and Jun Guo</i>	
Context-Guided Diffusion for Label Propagation on Graphs	2776
<i>Kwang In Kim, James Tompkin, Hanspeter Pfister, and Christian Theobalt</i>	
Learning to Rank Based on Subsequences	2785
<i>Basura Fernando, Efstratios Gavves, Damien Moselet, and Tinne Tuytelaars</i>	
Unsupervised Learning of Visual Representations Using Videos	2794
<i>Xiaolong Wang and Abhinav Gupta</i>	
A Nonparametric Bayesian Approach toward Stacked Convolutional Independent Component Analysis	2803
<i>Sotirios P. Chatzis and Dimitrios Kosmopoulos</i>	
Robust Principal Component Analysis on Graphs	2812
<i>Nauman Shahid, Vassilis Kalofolias, Xavier Bresson, Michael Bronstein, and Pierre Vandergheynst</i>	
Projection Bank: From High-Dimensional Data to Medium-Length Binary Codes	2821
<i>Li Liu, Mengyang Yu, and Ling Shao</i>	
Robust Optimization for Deep Regression	2830
<i>Vasileios Belagiannis, Christian Rupprecht, Gustavo Carneiro, and Nassir Navab</i>	
Multi-class Multi-annotator Active Learning with Robust Gaussian Process for Visual Recognition	2839
<i>Chengjiang Long and Gang Hua</i>	
Maximum-Margin Structured Learning with Deep Networks for 3D Human Pose Estimation	2848
<i>Sijin Li, Weichen Zhang, and Antoni B. Chan</i>	
An Exploration of Parameter Redundancy in Deep Networks with Circulant Projections	2857
<i>Yu Cheng, Felix X. Yu, Rogerio S. Feris, Sanjiv Kumar, Alok Choudhary, and Shi-Fu Chang</i>	
Additive Nearest Neighbor Feature Maps	2866
<i>Zhenzhen Wang, Xiao-Tong Yuan, Qingshan Liu, and Shuicheng Yan</i>	
Understanding Deep Features with Computer-Generated Imagery	2875
<i>Mathieu Aubry and Bryan C. Russell</i>	

Interpolation on the Manifold of K Component GMMs	2884
<i>Hyunwoo J. Kim, Nagesh Adluru, Monami Banerjee, Baba C. Vemuri, and Vikas Singh</i>	
Context-Aware CNNs for Person Head Detection	2893
<i>Tuan-Hung Vu, Anton Osokin, and Ivan Laptev</i>	
Mode-Seeking on Hypergraphs for Robust Geometric Model Fitting	2902
<i>Hanzi Wang, Guobao Xiao, Yan Yan, and David Suter</i>	
Highly-Expressive Spaces of Well-Behaved Transformations: Keeping it Simple	2911
<i>Oren Freifeld, Søren Hauberg, Kayhan Batmanghelich, and John W. Fisher III</i>	
Entropy-Based Latent Structured Output Prediction	2920
<i>Diane Bouchacourt, Sebastian Nowozin, and M. Pawan Kumar</i>	
Fast Orthogonal Projection Based on Kronecker Product	2929
<i>Xu Zhang, Felix X. Yu, Ruiqi Guo, Sanjiv Kumar, Shengjin Wang, and Shi-Fu Chang</i>	
PoseNet: A Convolutional Network for Real-Time 6-DOF Camera Relocalization	2938
<i>Alex Kendall, Matthew Grimes, and Roberto Cipolla</i>	
Predicting Multiple Structured Visual Interpretations	2947
<i>Debadeepta Dey, Varun Ramakrishna, Martial Hebert, and J. Andrew Bagnell</i>	
Look and Think Twice: Capturing Top-Down Visual Attention with Feedback Convolutional Neural Networks	2956
<i>Chunshui Cao, Xianming Liu, Yi Yang, Yinan Yu, Jiang Wang, Zilei Wang, Yongzhen Huang, Liang Wang, Chang Huang, Wei Xu, Deva Ramanan, and Thomas S. Huang</i>	
Matrix Backpropagation for Deep Networks with Structured Layers	2965
<i>Catalin Ionescu, Orestis Vantzos, and Cristian Sminchisescu</i>	
Introducing Geometry in Active Learning for Image Segmentation	2974
<i>Ksenia Konyushkova, Raphael Sznitman, and Pascal Fua</i>	
Joint Fine-Tuning in Deep Neural Networks for Facial Expression Recognition	2983
<i>Heechul Jung, Sihaeng Lee, Junho Yim, Sunjeong Park, and Junmo Kim</i>	
Direct Intrinsics: Learning Albedo-Shading Decomposition by Convolutional Regression	2992
<i>Takuya Narihira, Michael Maire, and Stella X. Yu</i>	
Face Flow	2993
<i>Patrick Snape, Anastasios Roussos, Yannis Panagakis, and Stefanos Zafeiriou</i>	
Discriminative Low-Rank Tracking	3002
<i>Yao Sui, Yafei Tang, and Li Zhang</i>	
SOWP: Spatially Ordered and Weighted Patch Descriptor for Visual Tracking	3011
<i>Han-UI Kim, Dae-Youn Lee, Jae-Young Sim, and Chang-Su Kim</i>	
Live Repetition Counting	3020
<i>Ofir Levy and Lior Wolf</i>	
Near-Online Multi-target Tracking with Aggregated Local Flow Descriptor	3029
<i>Wongun Choi</i>	
Multi-kernel Correlation Filter for Visual Tracking	3038
<i>Ming Tang and Jiayi Feng</i>	

Joint Probabilistic Data Association Revisited	3047
<i>Seyed Hamid Rezatofighi, Anton Milan, Zhen Zhang, Qinfeng Shi, Anthony Dick, and Ian Reid</i>	
Tracking-by-Segmentation with Online Gradient Boosting Decision Tree	3056
<i>Jeany Son, Ilchae Jung, Kayoung Park, and Bohyung Han</i>	
Exploring Causal Relationships in Visual Object Tracking	3065
<i>Karel Lebeda, Simon Hadfield, and Richard Bowden</i>	
Hierarchical Convolutional Features for Visual Tracking	3074
<i>Chao Ma, Jia-Bin Huang, Xiaokang Yang, and Ming-Hsuan Yang</i>	
Robust Non-rigid Motion Tracking and Surface Reconstruction Using L0 Regularization	3083
<i>Kaiwen Guo, Feng Xu, Yangang Wang, Yebin Liu, and Qionghai Dai</i>	
Online Object Tracking with Proposal Selection	3092
<i>Yang Hua, Karteek Alahari, and Cordelia Schmid</i>	
Understanding and Diagnosing Visual Tracking Systems	3101
<i>Naiyan Wang, Jianping Shi, Dit-Yan Yeung, and Jiaya Jia</i>	
Integrating Dashcam Views through Inter-Video Mapping	3110
<i>Hsin-I Chen, Yi-Ling Chen, Wei-Tse Lee, Fan Wang, and Bing-Yu Chen</i>	
Visual Tracking with Fully Convolutional Networks	3119
<i>Lijun Wang, Wanli Ouyang, Xiaogang Wang, and Huchuan Lu</i>	
Multiple Feature Fusion via Weighted Entropy for Visual Tracking	3128
<i>Lin Ma, Jiwen Lu, Jianjiang Feng, and Jie Zhou</i>	
Pedestrian Travel Time Estimation in Crowded Scenes	3137
<i>Shuai Yi, Hongsheng Li, and Xiaogang Wang</i>	
Unsupervised Synchrony Discovery in Human Interaction	3146
<i>Wen-Sheng Chu, Jiabei Zeng, Fernando De la Torre, Jeffrey F. Cohn, and Daniel S. Messinger</i>	
Efficient Video Segmentation Using Parametric Graph Partitioning	3155
<i>Chen-Ping Yu, Hieu Le, Gregory Zelinsky, and Dimitris Samaras</i>	
Learning to Track for Spatio-Temporal Action Localization	3164
<i>Philippe Weinzaepfel, Zaid Harchaoui, and Cordelia Schmid</i>	
Unsupervised Object Discovery and Tracking in Video Collections	3173
<i>Suha Kwak, Minsu Cho, Ivan Laptev, Jean Ponce, and Cordelia Schmid</i>	
Car that Knows Before You Do: Anticipating Maneuvers via Learning Temporal Driving Models	3182
<i>Ashesh Jain, Hema S. Koppula, Bharad Raghavan, Shane Soh, and Ashutosh Saxena</i>	
Activity Auto-Completion: Predicting Human Activities from Partial Videos	3191
<i>Zhen Xu, Laiyun Qing, and Jun Miao</i>	
Person Re-Identification with Correspondence Structure Learning	3200
<i>Yang Shen, Weiyao Lin, Junchi Yan, Mingliang Xu, Jianxin Wu, and Jingdong Wang</i>	
Adaptive Exponential Smoothing for Online Filtering of Pixel Prediction Maps	3209
<i>Kang Dang, Jiong Yang, and Junsong Yuan</i>	

P-CNN: Pose-Based CNN Features for Action Recognition	3218
<i>Guilhem Chéron, Ivan Laptev, and Cordelia Schmid</i>	
Fully Connected Object Proposals for Video Segmentation	3227
<i>Federico Perazzi, Oliver Wang, Markus Gross, and Alexander Sorkine-Hornung</i>	
Video Segmentation with Just a Few Strokes	3235
<i>Naveen Shankar Nagaraja, Frank R. Schmidt, and Thomas Brox</i>	
Actionness-Assisted Recognition of Actions	3244
<i>Ye Luo, Loong-Fah Cheong, and An Tran</i>	
COUNT Forest: CO-Voting Uncertain Number of Targets Using Random Forest for Crowd Density Estimation	3253
<i>Viet-Quoc Pham, Tatsuo Kozakaya, Osamu Yamaguchi, and Ryuzo Okada</i>	
Multi-cue Structure Preserving MRF for Unconstrained Video Segmentation	3262
<i>Sae-hoon Yi and Vladimir Pavlovic</i>	
Motion Trajectory Segmentation via Minimum Cost Multicuts	3271
<i>Margret Keuper, Bjoern Andres, and Thomas Brox</i>	
Action Localization in Videos through Context Walk	3280
<i>Khurram Soomro, Haroon Idrees, and Mubarak Shah</i>	
RGB-W: When Vision Meets Wireless	3289
<i>Alexandre Alahi, Albert Haque, and Li Fei-Fei</i>	
Action Detection by Implicit Intentional Motion Clustering	3298
<i>Wei Chen and Jason J. Corso</i>	
Simultaneous Foreground Detection and Classification with Hybrid Features	3307
<i>Jaemyun Kim, Adin Ramírez Rivera, Byungyong Ryu, and Oksam Chae</i>	

O-3C: Vision and People

Training a Feedback Loop for Hand Pose Estimation	3316
<i>Markus Oberweger, Paul Wohlhart, and Vincent Lepetit</i>	
Opening the Black Box: Hierarchical Sampling Optimization for Estimating Human Hand Pose	3325
<i>Danhang Tang, Jonathan Taylor, Pushmeet Kohli, Cem Keskin, Tae-Kyun Kim, and Jamie Shotton</i>	
Panoptic Studio: A Massively Multiview System for Social Motion Capture	3334
<i>Hanbyul Joo, Hao Liu, Lei Tan, Lin Gui, Bart Nabbe, Iain Matthews, Takeo Kanade, Shohei Nobuhara, and Yaser Sheikh</i>	
Where to Buy It: Matching Street Clothing Photos in Online Shops	3343
<i>M. Hadi Kiapour, Xufeng Han, Svetlana Lazebnik, Alexander C. Berg, and Tamara L. Berg</i>	
Multi-task Recurrent Neural Network for Immediacy Prediction	3352
<i>Xiao Chu, Wanli Ouyang, Wei Yang, and Xiaogang Wang</i>	
Learning Complexity-Aware Cascades for Deep Pedestrian Detection	3361
<i>Zhaowei Cai, Mohammad Saberian, and Nuno Vasconcelos</i>	

O-4A: Computational Photography and Image Enhancement

Polarized 3D: High-Quality Depth Sensing with Polarization Cues	3370
<i>Achuta Kadambi, Vage Taamazyan, Boxin Shi, and Ramesh Raskar</i>	
Airborne Three-Dimensional Cloud Tomography	3379
<i>Aviad Levis, Yoav Y. Schechner, Amit Aides, and Anthony B. Davis</i>	
Leave-One-Out Kernel Optimization for Shadow Detection	3388
<i>Tomás F. Yago Vicente, Minh Hoai, and Dimitris Samaras</i>	
Removing Rain from a Single Image via Discriminative Sparse Coding	3397
<i>Yu Luo, Yong Xu, and Hui Ji</i>	
Mutual-Structure for Joint Filtering	3406
<i>Xiaoyong Shen, Chao Zhou, Li Xu, and Jiaya Jia</i>	

P-4A: Computational Photography, Face and Gesture, and Vision for X

Photometric Stereo in a Scattering Medium	3415
<i>Zak Murez, Tali Treibitz, Ravi Ramamoorthi, and David Kriegman</i>	
Resolving Scale Ambiguity via XSlit Aspect Ratio Analysis	3424
<i>Wei Yang, Haiting Lin, Sing Bing Kang, and Jingyi Yu</i>	
Single-Shot Specular Surface Reconstruction with Gonio-Plenoptic Imaging	3433
<i>Lingfei Meng, Liyang Lu, Noah Bedard, and Kathrin Berkner</i>	
TransCut: Transparent Object Segmentation from a Light-Field Image	3442
<i>Yichao Xu, Hajime Nagahara, Atsushi Shimada, and Rin-ichiro Taniguchi</i>	
Depth Recovery from Light Field Using Focal Stack Symmetry	3451
<i>Haiting Lin, Can Chen, Sing Bing Kang, and Jingyi Yu</i>	
Depth Map Estimation and Colorization of Anaglyph Images Using Local Color Prior and Reverse Intensity Distribution	3460
<i>Williem, Ramesh Raskar, and In Kyu Park</i>	
Learning Data-Driven Reflectance Priors for Intrinsic Image Decomposition	3469
<i>Tinghui Zhou, Philipp Krähenbühl, and Alexei A. Efros</i>	
Photometric Stereo with Small Angular Variations	3478
<i>Jian Wang, Yasuyuki Matsushita, Boxin Shi, and Aswin C. Sankaranarayanan</i>	
Occlusion-Aware Depth Estimation Using Light-Field Cameras	3487
<i>Ting-Chun Wang, Alexei A. Efros, and Ravi Ramamoorthi</i>	
Oriented Light-Field Windows for Scene Flow	3496
<i>Pratul P. Srinivasan, Michael W. Tao, Ren Ng, and Ravi Ramamoorthi</i>	
Extended Depth of Field Catadioptric Imaging Using Focal Sweep	3505
<i>Ryunosuke Yokoya and Shree K. Nayar</i>	
Intrinsic Depth: Improving Depth Transfer with Intrinsic Images	3514
<i>Naejin Kong and Michael J. Black</i>	

Separating Fluorescent and Reflective Components by Using a Single Hyperspectral Image	3523
<i>Yinqiang Zheng, Ying Fu, Antony Lam, Imari Sato, and Yoichi Sato</i>	
Frequency-Based Environment Matting by Compressive Sensing	3532
<i>Yiming Qian, Minglun Gong, and Yee-Hong Yang</i>	
Complementary Sets of Shutter Sequences for Motion Deblurring	3541
<i>Hae-Gon Jeon, Joon-Young Lee, Yudeog Han, Seon Joo Kim, and In So Kweon</i>	
Hyperspectral Compressive Sensing Using Manifold-Structured Sparsity Prior	3550
<i>Lei Zhang, Wei Wei, Yanning Zhang, Fei Li, Chunhua Shen, and Qinfeng Shi</i>	
A Gaussian Process Latent Variable Model for BRDF Inference	3559
<i>Stamatios Georgoulis, Vincent Vanweddigen, Marc Proesmans, and Luc Van Gool</i>	
Active One-Shot Scan for Wide Depth Range Using a Light Field Projector Based on Coded Aperture	3568
<i>Hiroshi Kawasaki, Satoshi Ono, Yuki Horita, Yuki Shiba, Ryo Furukawa, and Shinsaku Hiura</i>	
Model-Based Tracking at 300Hz Using Raw Time-of-Flight Observations	3577
<i>Jan Stühmer, Sebastian Nowozin, Andrew Fitzgibbon, Richard Szeliski, Travis Perry, Sunil Acharya, Daniel Cremers, and Jamie Shotton</i>	
Hyperspectral Super-Resolution by Coupled Spectral Unmixing	3586
<i>Charis Lanaras, Emmanuel Baltsavias, and Konrad Schindler</i>	
Depth Selective Camera: A Direct, On-Chip, Programmable Technique for Depth Selectivity in Photography	3595
<i>Ryuichi Tadano, Adithya Kumar Pediredla, and Ashok Veeraraghavan</i>	
A Groupwise Multilinear Correspondence Optimization for 3D Faces	3604
<i>Timo Bolkart and Stefanie Wuhrer</i>	
Selective Encoding for Recognizing Unreliably Localized Faces	3613
<i>Ang Li, Vlad I. Morariu, and Larry S. Davis</i>	
Confidence Preserving Machine for Facial Action Unit Detection	3622
<i>Jiabei Zeng, Wen-Sheng Chu, Fernando De la Torre, Jeffrey F. Cohn, and Zhang Xiong</i>	
Learning Social Relation Traits from Face Images	3631
<i>Zhanpeng Zhang, Ping Luo, Chen-Change Loy, and Xiaoou Tang</i>	
Robust Heart Rate Measurement from Video Using Select Random Patches	3640
<i>Antony Lam and Yoshinori Kuno</i>	
Robust Model-Based 3D Head Pose Estimation	3649
<i>Gregory P. Meyer, Shalini Gupta, Iuri Frosio, Dikpal Reddy, and Jan Kautz</i>	
Robust Facial Landmark Detection Under Significant Head Poses and Occlusion	3658
<i>Yue Wu and Qiang Ji</i>	
Conditional Convolutional Neural Network for Modality-Aware Face Recognition	3667
<i>Chao Xiong, Xiaowei Zhao, Danhang Tang, Karlekar Jayashree, Shuicheng Yan, and Tae-Kyun Kim</i>	
From Facial Parts Responses to Face Detection: A Deep Learning Approach	3676
<i>Shuo Yang, Ping Luo, Chen-Change Loy, and Xiaoou Tang</i>	

Efficient PSD Constrained Asymmetric Metric Learning for Person Re-Identification	3685
<i>Shengcai Liao and Stan Z. Li</i>	
Pose-Invariant 3D Face Alignment	3694
<i>Amin Jourabloo and Xiaoming Liu</i>	
From Emotions to Action Units with Hidden and Semi-Hidden-Task Learning	3703
<i>Adria Ruiz, Joost Van de Weijer, and Xavier Binefa</i>	
Automated Facial Trait Judgment and Election Outcome Prediction: Social Dimensions of Face	3712
<i>Jungseock Joo, Francis F. Steen, and Song-Chun Zhu</i>	
Simultaneous Local Binary Feature Learning and Encoding for Face Recognition	3721
<i>Jiwen Lu, Venice Erin Liong, and Jie Zhou</i>	
Deep Learning Face Attributes in the Wild	3730
<i>Ziwei Liu, Ping Luo, Xiaogang Wang, and Xiaoou Tang</i>	
Multi-Task Learning with Low Rank Attribute Embedding for Person Re-Identification	3739
<i>Chi Su, Fan Yang, Shiliang Zhang, Qi Tian, Larry S. Davis, and Wen Gao</i>	
Regressing a 3D Face Shape from a Single Image	3748
<i>Sergey Tulyakov and Nicu Sebe</i>	
Rendering of Eyes for Eye-Shape Registration and Gaze Estimation	3756
<i>Erroll Wood, Tadas Baltrušaitis, Xucong Zhang, Yusuke Sugano, Peter Robinson, and Andreas Bulling</i>	
Multi-Scale Learning for Low-Resolution Person Re-Identification	3765
<i>Xiang Li, Wei-Shi Zheng, Xiaojuan Wang, Tao Xiang, and Shaogang Gong</i>	
Learning to Transfer: Transferring Latent Task Structures and Its Application to Person-Specific Facial Action Unit Detection	3774
<i>Timur Almaev, Brais Martinez, and Michel Valstar</i>	
Pairwise Conditional Random Forests for Facial Expression Recognition	3783
<i>Arnaud Dapogny, Kevin Bailly, and Séverine Dubuisson</i>	
Multi-conditional Latent Variable Model for Joint Facial Action Unit Detection	3792
<i>Stefanos Eleftheriadis, Ognjen Rudovic, and Maja Pantic</i>	
Leveraging Datasets with Varying Annotations for Face Alignment via Deep Regression Network	3801
<i>Jie Zhang, Meina Kan, Shiguang Shan, and Xilin Chen</i>	
A Spatio-Temporal Appearance Representation for Video-Based Pedestrian Re-Identification	3810
<i>Kan Liu, Bingpeng Ma, Wei Zhang, and Rui Huang</i>	
Two Birds, One Stone: Jointly Learning Binary Code for Large-Scale Face Image Retrieval and Attributes Prediction	3819
<i>Yan Li, Ruiping Wang, Haomiao Liu, Huajie Jiang, Shiguang Shan, and Xilin Chen</i>	
An Accurate Iris Segmentation Framework Under Relaxed Imaging Constraints Using Total Variation Model	3828
<i>Zijing Zhao and Ajay Kumar</i>	

Discriminative Pose-Free Descriptors for Face and Object Matching	3837
<i>Soubhik Sanyal, Sivaram Prasad Mudunuri, and Soma Biswas</i>	
Bi-Shifting Auto-Encoder for Unsupervised Domain Adaptation	3846
<i>Meina Kan, Shiguang Shan, and Xilin Chen</i>	
Regressive Tree Structured Model for Facial Landmark Localization	3855
<i>Gee-Sern (Jison) Hsu, Kai-Hsiang Chang, and Shih-Chieh Huang</i>	
Person Recognition in Personal Photo Collections	3862
<i>Seong Joon Oh, Rodrigo Benenson, Mario Fritz, and Bernt Schiele</i>	
Robust Statistical Face Frontalization	3871
<i>Christos Sagonas, Yannis Panagakis, Stefanos Zafeiriou, and Maja Pantic</i>	
PIEFA: Personalized Incremental and Ensemble Face Alignment	3880
<i>Xi Peng, Shaoting Zhang, Yu Yang, and Dimitris N. Metaxas</i>	
Understanding Everyday Hands in Action from RGB-D Images	3889
<i>Grégory Rogez, James S. Supančič III, and Deva Ramanan</i>	
Example-Based Modeling of Facial Texture from Deficient Data	3898
<i>Arnaud Dessein, William A. P. Smith, Richard C. Wilson, and Edwin R. Hancock</i>	
Learning to Predict Saliency on Face Images	3907
<i>Mai Xu, Yun Ren, and Zulin Wang</i>	
Group Membership Prediction	3916
<i>Ziming Zhang, Yuting Chen, and Venkatesh Saligrama</i>	
Extraction of Virtual Baselines from Distorted Document Images Using Curvilinear Projection	3925
<i>Gaofeng Meng, Zuming Huang, Yonghong Song, Shiming Xiang, and Chunhong Pan</i>	
Robust RGB-D Odometry Using Point and Line Features	3934
<i>Yan Lu and Dezhen Song</i>	
Learning a Discriminative Model for the Perception of Realism in Composite Images	3943
<i>Jun-Yan Zhu, Philipp Krähenbühl, Eli Shechtman, and Alexei A. Efros</i>	
What Makes Tom Hanks Look Like Tom Hanks	3952
<i>Supasorn Suwajanakorn, Steven M. Seitz, and Ira Kemelmacher-Shlizerman</i>	
Wide-Area Image Geolocalization with Aerial Reference Imagery	3961
<i>Scott Workman, Richard Souvenir, and Nathan Jacobs</i>	
Personalized Age Progression with Aging Dictionary	3970
<i>Xiangbo Shu, Jinhui Tang, Hanjiang Lai, Luoqi Liu, and Shuicheng Yan</i>	
FaceDirector: Continuous Control of Facial Performance in Video	3979
<i>Charles Malleson, Jean-Charles Bazin, Oliver Wang, Derek Bradley, Thabo Beeler, Adrian Hilton, and Alexander Sorkine-Hornung</i>	
Synthesizing Illumination Mosaics from Internet Photo-Collections	3988
<i>Dinghuang Ji, Enrique Dunn, and Jan-Michael Frahm</i>	
Hot or Not: Exploring Correlations between Appearance and Temperature	3997
<i>Daniel Glasner, Pascal Fua, Todd Zickler, and Lihi Zelnik-Manor</i>	

O-4B: Motion and Correspondence

SPM-BP: Sped-Up PatchMatch Belief Propagation for Continuous MRFs	4006
<i>Yu Li, Dongbo Min, Michael S. Brown, Minh N. Do, and Jiangbo Lu</i>	
Flow Fields: Dense Correspondence Fields for Highly Accurate Large Displacement Optical Flow Estimation	4015
<i>Christian Bailer, Bertram Taetz, and Didier Stricker</i>	
Dense Semantic Correspondence Where Every Pixel is a Classifier	4024
<i>Hilton Bristow, Jack Valmadre, and Simon Lucey</i>	
Multi-image Matching via Fast Alternating Minimization	4032
<i>Xiaowei Zhou, Menglong Zhu, and Kostas Daniilidis</i>	

P-4B: Statistical Methods and Learning, Motion and Tracking, and Video Analysis II

Differential Recurrent Neural Networks for Action Recognition	4041
<i>Vivek Veeriah, Naifan Zhuang, and Guo-Jun Qi</i>	
Similarity Gaussian Process Latent Variable Model for Multi-modal Data Analysis	4050
<i>Guoli Song, Shuhui Wang, Qingming Huang, and Qi Tian</i>	
Learning Ensembles of Potential Functions for Structured Prediction with Latent Variables	4059
<i>Hossein Hajimirsadeghi and Greg Mori</i>	
Simultaneous Deep Transfer Across Domains and Tasks	4068
<i>Eric Tzeng, Judy Hoffman, Trevor Darrell, and Kate Saenko</i>	
Low Dimensional Explicit Feature Maps	4077
<i>Ondřej Chum</i>	
Unsupervised Learning of Spatiotemporally Coherent Metrics	4086
<i>Ross Goroshin, Joan Bruna, Jonathan Tompson, David Eigen, and Yann LeCun</i>	
Multi-label Cross-Modal Retrieval	4094
<i>Viresh Ranjan, Nikhil Rasiwasia, and C. V. Jawahar</i>	
Improving Ferns Ensembles by Sparsifying and Quantising Posterior Probabilities	4103
<i>Antonio L. Rodriguez and Vitor Sequeira</i>	
Beyond Gauss: Image-Set Matching on the Riemannian Manifold of PDFs	4112
<i>Mehrtash Harandi, Mathieu Salzmann, and Mahsa Baktashmotlagh</i>	
Unsupervised Domain Adaptation with Imbalanced Cross-Domain Data	4121
<i>Tzu Ming Harry Hsu, Wei Yu Chen, Cheng-An Hou, Yao-Hung Hubert Tsai, Yi-Ren Yeh, and Yu-Chiang Frank Wang</i>	
Secrets of Matrix Factorization: Approximations, Numerics, Manifold Optimization and Random Restarts	4130
<i>Je Hyeong Hong and Andrew Fitzgibbon</i>	
Geometry-Aware Deep Transform	4139
<i>Jiaji Huang, Qiang Qiu, Robert Calderbank, and Guillermo Sapiro</i>	

Learning Binary Codes for Maximum Inner Product Search	4148
<i>Fumin Shen, Wei Liu, Shaoting Zhang, Yang Yang, and Heng Tao Shen</i>	
ML-MG: Multi-label Learning with Missing Labels Using a Mixed Graph	4157
<i>Baoyuan Wu, Siwei Lyu, and Bernard Ghanem</i>	
Zero-Shot Learning via Semantic Similarity Embedding	4166
<i>Ziming Zhang and Venkatesh Saligrama</i>	
Bayesian Model Adaptation for Crowd Counts	4175
<i>Bo Liu and Nuno Vasconcelos</i>	
An NMF Perspective on Binary Hashing	4184
<i>Lopamudra Mukherjee, Sathya N. Ravi, Vamsi K. Ithapu, Tyler Holmes, and Vikas Singh</i>	
Multi-view Domain Generalization for Visual Recognition	4193
<i>Li Niu, Wen Li, and Dong Xu</i>	
Infinite Feature Selection	4202
<i>Giorgio Roffo, Simone Melzi, and Marco Cristani</i>	
Semi-Supervised Zero-Shot Classification with Label Representation Learning	4211
<i>Xin Li, Yuhong Guo, and Dale Schuurmans</i>	
A Supervised Low-Rank Method for Learning Invariant Subspaces	4220
<i>Farzad Siyahjani, Ranya Almohsen, Sinan Sabri, and Gianfranco Doretto</i>	
Recursive Fréchet Mean Computation on the Grassmannian and Its Applications to Computer Vision	4229
<i>Rudrasis Chakraborty and Baba C. Vemuri</i>	
Multi-view Subspace Clustering	4238
<i>Hongchang Gao, Feiping Nie, Xuelong Li, and Heng Huang</i>	
Predicting Deep Zero-Shot Convolutional Neural Networks Using Textual Descriptions	4247
<i>Jimmy Lei Ba, Kevin Swersky, Sanja Fidler, and Ruslan Salakhutdinov</i>	
Structured Feature Selection	4256
<i>Tian Gao, Ziheng Wang, and Qiang Ji</i>	
Conditional High-Order Boltzmann Machine: A Supervised Learning Model for Relation Learning	4265
<i>Yan Huang, Wei Wang, and Liang Wang</i>	
Learning Image and User Features for Recommendation in Social Networks	4274
<i>Xue Geng, Hanwang Zhang, Jingwen Bian, and Tat-Seng Chua</i>	
Dual-Feature Warping-Based Motion Model Estimation	4283
<i>Shiwei Li, Lu Yuan, Jian Sun, and Long Quan</i>	
An Adaptive Data Representation for Robust Point-Set Registration and Merging	4292
<i>Dylan Campbell and Lars Petersson</i>	
Local Subspace Collaborative Tracking	4301
<i>Lin Ma, Xiaoqin Zhang, Weiming Hu, Junliang Xing, Jiwen Lu, and Jie Zhou</i>	
Learning Spatially Regularized Correlation Filters for Visual Tracking	4310
<i>Martin Danelljan, Gustav Häger, Fahad Shahbaz Khan, and Michael Felsberg</i>	

SpeDo: 6 DOF Ego-Motion Sensor Using Speckle Defocus Imaging	4319
<i>Kensei Jo, Mohit Gupta, and Shree K. Nayar</i>	
Unsupervised Trajectory Clustering via Adaptive Multi-kernel-Based Shrinkage	4328
<i>Hongteng Xu, Yang Zhou, Weiyao Lin, and Hongyuan Zha</i>	
TRIC-track: Tracking by Regression with Incrementally Learned Cascades	4337
<i>Xiaomeng Wang, Michel Valstar, Brais Martinez, Muhammad Haris Khan, and Tony Pridmore</i>	
Recurrent Network Models for Human Dynamics	4346
<i>Katerina Fragkiadaki, Sergey Levine, Panna Felsen, and Jitendra Malik</i>	
Contour Flow: Middle-Level Motion Estimation by Combining Motion Segmentation and Contour Alignment	4355
<i>Huijun Di, Qingxuan Shi, Feng Lv, Ming Qin, and Yao Lu</i>	
FollowMe: Efficient Online Min-Cost Flow Tracking with Bounded Memory and Computation	4364
<i>Philip Lenz, Andreas Geiger, and Raquel Urtasun</i>	
Learning to Divide and Conquer for Online Multi-target Tracking	4373
<i>Francesco Solera, Simone Calderara, and Rita Cucchiara</i>	
Minimizing Human Effort in Interactive Tracking by Incremental Learning of Model Parameters	4382
<i>Arridhana Ciptadi and James M. Rehg</i>	
A Novel Representation of Parts for Accurate 3D Object Detection and Tracking in Monocular Images	4391
<i>Alberto Crivellaro, Mahdi Rad, Yannick Verdie, Kwang Moo Yi, Pascal Fua, and Vincent Lepetit</i>	
Linearization to Nonlinear Learning for Visual Tracking	4400
<i>Bo Ma, Hongwei Hu, Jianbing Shen, Yuping Zhang, and Fatih Porikli</i>	
Self-Occlusions and Disocclusions in Causal Video Object Segmentation	4408
<i>Yanchao Yang, Ganesh Sundaramoorthi, and Stefano Soatto</i>	
Large Displacement 3D Scene Flow with Occlusion Reasoning	4417
<i>Andrei Zanfir and Cristian Sminchisescu</i>	
Co-Interest Person Detection from Multiple Wearable Camera Videos	4426
<i>Yuewei Lin, Kareem Abdelfatah, Youjie Zhou, Xiaochuan Fan, Hongkai Yu, Hui Qian, and Song Wang</i>	
Sparse Dynamic 3D Reconstruction from Unsynchronized Videos	4435
<i>Enliang Zheng, Dinghuang Ji, Enrique Dunn, and Jan-Michael Frahm</i>	
Category-Blind Human Action Recognition: A Practical Recognition System	4444
<i>Wenbo Li, Longyin Wen, Mooi Choo Chuah, and Siwei Lyu</i>	
Temporal Subspace Clustering for Human Motion Segmentation	4453
<i>Sheng Li, Kang Li, and Yun Fu</i>	
Weakly-Supervised Alignment of Video with Text	4462
<i>P. Bojanowski, R. Lajugie, E. Grave, F. Bach, I. Laptev, J. Ponce, and C. Schmid</i>	

Learning Temporal Embeddings for Complex Video Analysis	4471
<i>Vignesh Ramanathan, Kevin Tang, Greg Mori, and Li Fei-Fei</i>	
Unsupervised Semantic Parsing of Video Collections	4480
<i>Ozan Sener, Amir R. Zamir, Silvio Savarese, and Ashutosh Saxena</i>	
Learning Spatiotemporal Features with 3D Convolutional Networks	4489
<i>Du Tran, Lubomir Bourdev, Rob Fergus, Lorenzo Torresani, and Manohar Paluri</i>	
Temporal Perception and Prediction in Ego-Centric Video	4498
<i>Yipin Zhou and Tamara L. Berg</i>	
Describing Videos by Exploiting Temporal Structure	4507
<i>Li Yao, Atousa Torabi, Kyunghyun Cho, Nicolas Ballas, Christopher Pal, Hugo Larochelle, and Aaron Courville</i>	
Person Re-Identification with Discriminatively Trained Viewpoint Invariant Dictionaries	4516
<i>Srikrishna Karanam, Yang Li, and Richard J. Radke</i>	
Storyline Representation of Egocentric Videos with an Applications to Story-Based Search	4525
<i>Bo Xiong, Gunhee Kim, and Leonid Sigal</i>	
Sequence to Sequence — Video to Text	4534
<i>Subhashini Venugopalan, Marcus Rohrbach, Jeffrey Donahue, Raymond Mooney, Trevor Darrell, and Kate Saenko</i>	
Context Aware Active Learning of Activity Recognition Models	4543
<i>Mahmudul Hasan and Amit K. Roy-Chowdhury</i>	
Action Recognition by Hierarchical Mid-Level Action Elements	4552
<i>Tian Lan, Yuke Zhu, Amir Roshan Zamir, and Silvio Savarese</i>	
Selecting Relevant Web Trained Concepts for Automated Event Retrieval	4561
<i>Bharat Singh, Xintong Han, Zhe Wu, Vlad I. Morariu, and Larry S. Davis</i>	
Beyond Covariance: Feature Representation with Nonlinear Kernel Matrices	4570
<i>Lei Wang, Jianjia Zhang, Luping Zhou, Chang Tang, and Wanqing Li</i>	
Multiresolution Hierarchy Co-Clustering for Semantic Segmentation in Sequences with Small Variations	4579
<i>David Varas, Mónica Alfaro, and Ferran Marques</i>	
Objects2action: Classifying and Localizing Actions without Any Video Example	4588
<i>Mihir Jain, Jan C. van Gemert, Thomas Mensink, and Cees G. M. Snoek</i>	
Human Action Recognition Using Factorized Spatio-Temporal Convolutional Networks	4597
<i>Lin Sun, Kui Jia, Dit-Yan Yeung, and Bertram E. Shi</i>	
Bayesian Non-parametric Inference for Manifold Based MoCap Representation	4606
<i>Fabrizio Natola, Valsamis Ntouskos, Marta Sanzari, and Fiora Pirri</i>	
Semantic Video Entity Linking Based on Visual Content and Metadata	4615
<i>Yuncheng Li, Xitong Yang, and Jiebo Luo</i>	
Love Thy Neighbors: Image Annotation by Exploiting Image Metadata	4624
<i>Justin Johnson, Lamberto Ballan, and Li Fei-Fei</i>	

Unsupervised Extraction of Video Highlights via Robust Recurrent Auto-Encoders	4633
<i>Huan Yang, Baoyuan Wang, Stephen Lin, David Wipf, Minyi Guo, and Baining Guo</i>	
Learning Visual Clothing Style with Heterogeneous Dyadic Co-Occurrences	4642
<i>Andreas Veit, Balazs Kovacs, Sean Bell, Julian McAuley, Kavita Bala, and Serge Belongie</i>	
Text Flow: A Unified Text Detection System in Natural Scene Images	4651
<i>Shangxuan Tian, Yifeng Pan, Chang Huang, Shijian Lu, Kai Yu, and Chew Lim Tan</i>	
 O-4C: Video: Actions, Surveillance and Tracking	
Uncovering Interactions and Interactors: Joint Estimation of Head, Body Orientation and F-Formations from Surveillance Videos	4660
<i>Elisa Ricci, Jagannadan Varadarajan, Ramanathan Subramanian, Samuel Rota Bulò, Narendra Ahuja, and Oswald Lenz</i>	
Generating Notifications for Missing Actions: Don't Forget to Turn the Lights Off!	4669
<i>Bilge Soran, Ali Farhadi, and Linda Shapiro</i>	
Partial Person Re-Identification	4678
<i>Wei-Shi Zheng, Xiang Li, Tao Xiang, Shengcai Liao, Jianhuang Lai, and Shaogang Gong</i>	
Shape Interaction Matrix Revisited and Robustified: Efficient Subspace Clustering with Corrupted and Incomplete Data	4687
<i>Pan Ji, Mathieu Salzmann, and Hongdong Li</i>	
Multiple Hypothesis Tracking Revisited	4696
<i>Chanho Kim, Fuxin Li, Arridhana Ciptadi, and James M. Rehg</i>	
Learning to Track: Online Multi-object Tracking by Decision Making	4705
<i>Yu Xiang, Alexandre Alahi, and Silvio Savarese</i>	
 Author Index	 4714